

# Performance Optimization - Cheat Sheet

---

## Amdahl's Law

---

[https://en.wikipedia.org/wiki/Amdahl%27s\\_law](https://en.wikipedia.org/wiki/Amdahl%27s_law)

$$S_{\text{latency}} = \frac{1}{1 - p + \frac{p}{s}}$$

Where  $S_{\text{latency}}$  is the theoretical speedup,  $s$  is the speedup of the part with improved system resources and  $p$  is the percentage of execution time that benefits from the improved resources.

Describes the theoretical speedup in latency of a task with a fixed workload when the system resources are improved. Usually used to describe speedup if a task is run on multiple processors.

One can use Amdahl's law to see if an upgrade is worth it or not.

### Example

By exchanging the disks in a server, one hopes to improve the overall performance. 30% of the system's performance relies on the disks. The new disks are twice as fast as the previous disks.

$$p = 0.3, s = 2$$
$$S_{\text{latency}} = \frac{1}{1 - 0.3 + \frac{0.3}{2}} \approx 1.18$$

Amdahl's law states that the latency of the system will be improved by 18%.

## Types of workloads

---

### Online Transaction Processing (OLTP)

This tasks revolve around interactive user activities. Generally OLTP includes measurements of a user's *think time*. The goal is usually to have a small response time.

The performance metric for OLTP is defined as the supported users  $N$ :

$$N_{\text{concurrent}} = N_{\text{active}} \left( \frac{\text{Response time}}{\text{Response time} + \text{think time}} \right)$$

To calculate the supported concurrent users one can use *real* workload (collected online) or *synthetic* workload. It is composed based on hypothetical use cases - usually described by stochastic processes.

### Batch jobs

Specifies the amount of work to be processed without the necessity of intervention (non-interactive).

The performance metric for batch jobs is *throughput* - the number of tasks completed within a given time period.

Calculating the throughput of batch jobs is relatively easy - assume that the system is always at maximum load.

## Queuing theory (QT)

---

Queuing theory is generic for computing and optimizing the efficiency of any system that achieves its objectives by consuming multiple resources. In customer call centers, computer systems etc. there's a limited amount of resources that should be optimally used to achieve maximum possible efficiency.

Provides a way to unify all metrics with quantitative performance laws.

### A metaphor - visiting a banking center.

- **Server:** banking center fulfilling the customer's service requests
- **Customer:** the initiator of service requests
- **Wait time:** the time duration a customer has to spend in line
- **Service time:** the time duration from when the teller starts the service a customer to when the customer is leaving the teller and the next customer is called in for a service
- **Arrival rate:** the rate at which customers arrive for service
- **Service rate:** the rate at which customers are serviced
- **Utilization:** the portion of a teller's time actually serving customers rather than idling
- **Queue length:** the total number of customer waiting or being serviced
- **Response time:** how long customers are in the system overall, on average. Usually called *latency*. Equal to wait time + service time
- **Residence time:** same as response time but when the teller handles multiple transactions per client
- **Throughput:** the number of users served per second
- **Demand:** number of visits \* service time

## RAID

---

RAID is a technology to set up an array of disks for different purposes. For example, RAID 1 uses one or more mirrored disks to achieve redundancy. RAID 0 uses two or more disks and write across them, a method sometimes called *striping*. RAID 5 uses at least three disks, two for RAID 0 and one for parity. RAID 5 is able to recover errors that may occur in one of the other disks.

## Memory Hierarchy

---

A computer has the following memory types:

1. Processor registers - the fastest type of memory around, very small
2. Cache Level 0-4. Ranges from a few KB with hundreds of GB/s access to over 128 MB with tens

of GB/s

3. Primary memory. Several GB, at least 10 GB/s access
4. Secondary memory. Multiple TB, hundreds of MB/s to some GB/s

## Server Oriented Architecture (SOA)

---

A reusable atomic service that does something independently or in a federated fashion - programming is about using services.

The user doesn't care about the *black box* - it sees the service.

Web Service Description Language (WSDL) is an XML file that describes the web service.

Writing a new application is not coding from scratch - it takes a whole bunch of ready-made "black-box" services.

## Optimization and tuning

---

Optimization is identifying and eliminating internal inefficient designs and implementations (code).

Tuning is establishing the optimal setting for every possible external configuration parameter (configuration).

## Little's law

Throughput, response time and queuing length completely quantifies the performance of a queuing system.

$$N_i = X_i R_i$$

$N$  is the number of customers waiting or receiving service,  $X_i$  is the throughput and  $R_i$  the response time.

## Hyper-threading

---

**Dual processor systems** work with separated sockets (physical processors).

**Hyper-threading** processors have shared ALUs and cache. The only separated parts are the logical processors and registers.

There is a need to synchronize logical CPUs. Single-threaded applications can actually slow down single-threaded applications. It benefits multi-threaded applications.

Not the same as two physical processors.

The effectiveness of hyper-threading depends on how busy the SUT without hyper-threading is under the intended load.

The OS sees a physical core and a logical core (from hyper-threading) appears as the same unit.

# Multicore

---

Multicore is a lot more powerful since a dual-core processor is closer to two processors than a single-core hyper-threading processor.

These cores share nothing above the L2 cache. Both single and multi-core applications can benefit from multicore.

Increases front-side-bus speed (CPU and memory controller) and size of L2/L3 cache.

More cores will help for CPU-bound workloads and when computation can be done in parallel.